

# Continuous Control with Deep Reinforcement Learning

Timothy P. Lillicrap, et al. · 14 p · 2015

By Flavio Schneider · November 2019 · **ETH** zürich



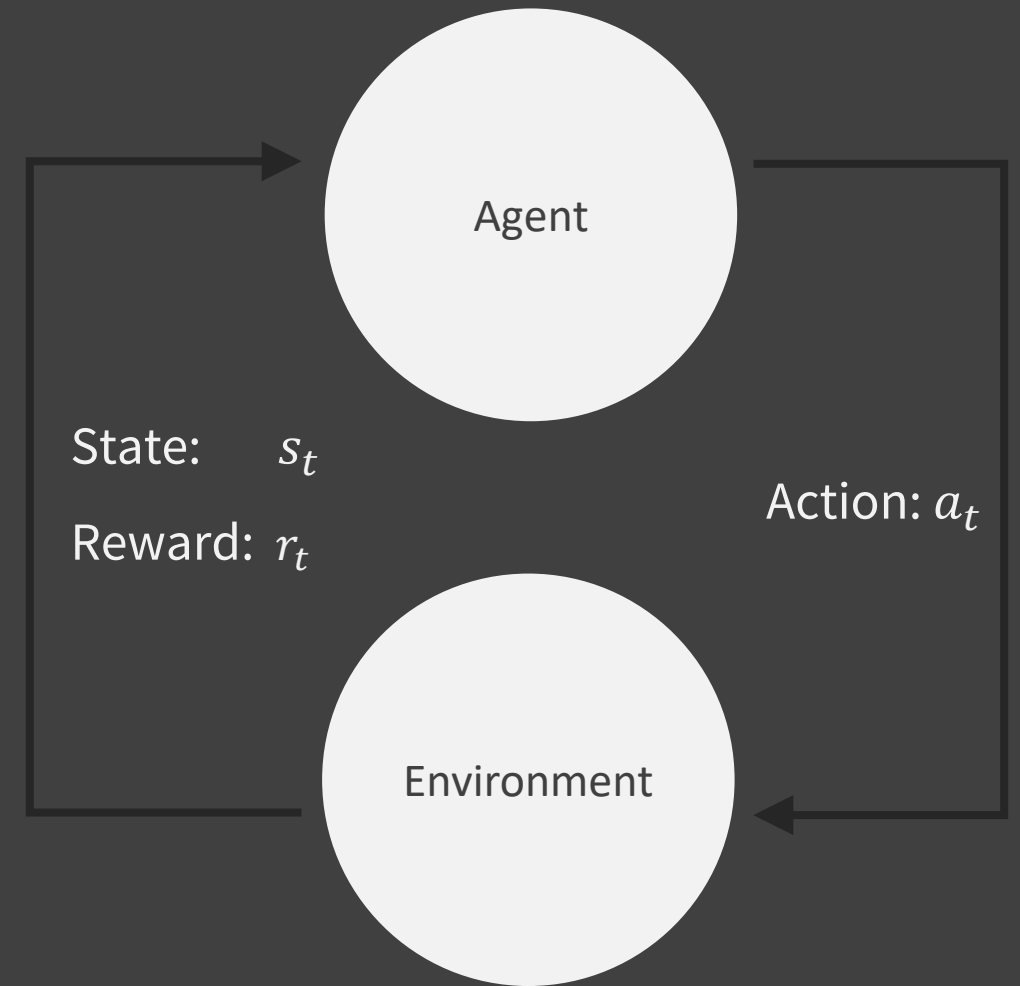
# Index

- What is reinforcement learning (RL)?
- How can we describe RL formally?
- How does a RL *algorithm* look like?

# What is Reinforcement Learning?

- One of the three basic paradigms of ML
- Learns intelligent behavior from reward
- Many applications
- Exploitation and exploration

Agent-Environment loop



# What is Reinforcement Learning?

State · Action · Reward · Policy · Goal

## Pendulum Example



# What is Reinforcement Learning?

State · Action · Reward · Policy · Goal

State:  $s_t \in S = \mathbb{R}^{n+1}$  describes to the agent the environment completely at time  $t$ .

State Vector

$$s_t = \begin{bmatrix} \text{Pendulum angle} \\ \text{Pendulum speed} \end{bmatrix}$$

$t = 0$




...

$t = 5$



...

$t = 10$



...

$$s_0 = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$s_5 = \begin{bmatrix} 20 \\ -5 \end{bmatrix}$$

$$s_{10} = \begin{bmatrix} -20 \\ 4 \end{bmatrix}$$

# What is Reinforcement Learning?

State · Action · Reward · Policy · Goal

Action:  $a_t \in A$  represents what the agent is going to do at time  $t$ .

Discrete Action Space

$$a_t \in A = (Left, Right) = (a^{(0)}, a^{(1)})$$



Continuous Action Space

$$a_t \in A = [-1, 1]$$



# What is Reinforcement Learning?

State · Action · Reward · Policy · Goal

Reward:  $r_t = R(s_t)$  or  $r_t = R(s_t, a_t)$  is a score that tells the agent how good was the last move  $a_t$  in state  $s_t$ .

Return:  $R(\tau) = \sum_{t=0}^T \gamma^t r_t$  where  $\tau = (s_0, a_0, s_1, a_1, \dots)$  is a sequence of state/action pairs and  $\gamma \in (0,1)$ .

Reward Function


$$r_t = R(s_t) = \begin{cases} 1 & \text{if } s_t[0] == 180 \\ 0 & \text{otherwise} \end{cases}$$

$t = 0$



$r_0 = 0$

$t = 100$



...

$r_{100} = 1$

# What is Reinforcement Learning?

State · Action · Reward · Policy · Goal

Policy:  $\mu(s_t) = a_t$  is the brain of the agent, a function that maps states to actions.



$$s_5 = \begin{bmatrix} 20 \\ -5 \end{bmatrix}$$



$$\mu\left(\begin{bmatrix} 20 \\ -5 \end{bmatrix}\right) = a_5^{(0)} = \textit{Left}$$



$$s_5 = \begin{bmatrix} 20 \\ -5 \end{bmatrix}$$



$$\mu\left(\begin{bmatrix} 20 \\ -5 \end{bmatrix}\right) = a_5^{(i)} = 0.83$$



# What is Reinforcement Learning?

State · Action · Reward · Policy · Goal

Trajectory Probability:

$$P(\tau \mid \mu) := \prod_{t=0}^{T-1} P(s_{t+1} \mid s_t, a_t) P(\mu(s_t) = a_t)$$

Expected Return:


$$\mathbb{E}_{a_t=\mu(s_t)}[R(\tau)] = \int_{\tau} P(\tau \mid \mu) R(\tau)$$

Goal · Find Optimal Policy:

$$\mu^* := \operatorname{argmax}_{\mu} \mathbb{E}_{a_t=\mu(s_t)}[R(\tau)]$$

Trajectory

$$\tau = (s_0, a_0^{(1)}, s_1, a_1^{(0)}, \dots)$$



# Important Functions

Quality · Bellmann

# Important Functions

Quality · Bellmann

Quality Function:

$$\mathbb{E}_{a_t=\mu(s_t)}[R(\tau)]$$


$$Q^\mu(s, a) := \mathbb{E}_{s_0=s, a_0=a, a_t=\mu(s_t)}[R(\tau)]$$

Optimal Quality Function:

$$Q^*(s, a) := \max_{\mu} Q^\mu(s, a)$$

Optimal Action:

$$\mu^*(s) = \underset{a}{\operatorname{argmax}} Q^*(s, a) = a^*$$



# Important Functions

Quality · Bellman

Bellman Equation:

$$Q^*(s, a) = \mathbb{E} \left[ R(s, a) + \gamma \max_{a'} Q^*(s', a') \right]$$

# Algorithms Evolution

DQN · DPG · DDPG


Question:

How can we solve any (discrete and then continuous)  
reinforcement learning problem?


Idea:

Approximate  $Q^*(s, a)$  using a neural network.

“Critic” Network ·  $Q^\mu(s, a \mid \theta^Q)$ :



Pendulum




# Algorithms Evolution

DQN · DPG · DDPG

Deep Q Network

1. Initialize randomly  $Q^\mu(s, a \mid \theta^Q)$
2. Get initial state  $s$
3. Repeat:
  - a.  $a = \underset{a'}{\operatorname{argmax}} Q^\mu(s, a' \mid \theta^Q)$
  - b.  $a = a^{(i)}$  for  $i \in \{0, \dots, m\}$  random with probability  $\epsilon$
  - c. Execute  $a$  and observe  $r = R(s, a)$  and  $s'$
  - d.  $Q^T(s, a, r, s') := r + \gamma \max_{a'} Q^\mu(s', a' \mid \theta^Q)$
  - e.  $\theta^Q \leftarrow \theta^Q - \alpha \nabla_{\theta^Q} \left[ \left( Q^T(s, a, r, s') - Q^\mu(s, a \mid \theta^Q) \right)^2 \right]$
  - f.  $s \leftarrow s'$

“Critic” Network ·  $Q^\mu(s, a \mid \theta^Q)$ :



Bellman Equation

$$Q^*(s, a) = \mathbb{E} \left[ R(s, a) + \gamma \max_{a'} Q^*(s', a') \right]$$


# Algorithms Evolution

DQN · DPG · DDPG

Deep Q Network

1. Initialize randomly  $Q^\mu(s, a \mid \theta^Q)$
2. Get initial state  $s$
3. Repeat:
  - a.  $a = \underset{a'}{\operatorname{argmax}} Q^\mu(s, a' \mid \theta^Q)$
  - b.  $a = a^{(i)}$  for  $i \in \{0, \dots, m\}$  random with probability  $\epsilon$
  - c. Execute  $a$  and observe  $r = R(s, a)$  and  $s'$
  - d.  $Q^T(s, a, r, s') := r + \gamma \max_{a'} Q^\mu(s', a' \mid \theta^Q)$
  - e.  $\theta^Q \leftarrow \theta^Q - \alpha \nabla_{\theta^Q} \left[ \left( Q^T(s, a, r, s') - Q^\mu(s, a \mid \theta^Q) \right)^2 \right]$
  - f.  $s \leftarrow s'$

“Critic” Network ·  $Q^\mu(s, a \mid \theta^Q)$ :



Bellman Equation



$$Q^*(s, a) = \mathbb{E} \left[ R(s, a) + \gamma \max_{a'} Q^*(s', a') \right]$$

# Algorithms Evolution


DQN · DPG · DDPG

Deep Q Network


## Code Structure

1. Initialize randomly  $Q^\mu(s, a | \theta^Q)$  and  $Q^{\mu'}(s, a | \theta^{Q'})$
2. Initialize replay buffer  $\mathcal{R} = S \times A \times R \times S$
3. Get initial state  $s$
4. Repeat:
  - a.  · Sample
  - b.  · Train

Target “Critic” Network ·  $Q^{\mu'}(s, a | \theta^{Q'})$  :



Replay Buffer ·  $\mathcal{R}$  :





# Algorithms Evolution


DQN · DPG · DDPG

Deep Q Network


■ · Sample

- i.  $a = \underset{a'}{\operatorname{argmax}} Q^\mu(s, a' \mid \theta^Q)$
- ii.  $a = a^{(i)}$  for  $i \in \{0, \dots, m\}$  random with probability  $\epsilon$
- iii. Execute  $a$  and observe  $r$  and  $s'$
- iv. Store  $\mathcal{R} \leftarrow \mathcal{R} \cup (s, a, r, s')$

Target “Critic” Network ·  $Q^{\mu'}(s, a \mid \theta^{Q'})$  :



Replay Buffer ·  $\mathcal{R}$  :



# Algorithms Evolution

DQN · DPG · DDPG

Deep Q Network

■ · Train

i. Sample random batch  $\mathcal{B} \subseteq \mathcal{R}$


ii. For each  $(s, a, r, s') \in \mathcal{B}$

$$I. \quad Q^T(s, a, r, s') := r + \gamma \max_{a'} Q^{\mu'}(s, a' | \theta^{Q'})$$


$$II. \quad \theta^Q \leftarrow \theta^Q - \alpha \nabla_{\theta^Q} \left[ \left( Q^T(s, a, r, s') - Q^{\mu}(s, a | \theta^Q) \right)^2 \right]$$

iii. Every  $C$  steps reset  $\theta^{Q'} \leftarrow \theta^Q$

Target “Critic” Network ·  $Q^{\mu'}(s, a | \theta^{Q'})$  :



Replay Buffer ·  $\mathcal{R}$  :



# Algorithms Evolution

DQN · DPG · DDPG

Deep Q Network

■ · Train

i. Sample random batch  $\mathcal{B} \subseteq \mathcal{R}$


ii. For each  $(s, a, r, s') \in \mathcal{B}$

$$I. \quad Q^T(s, a, r, s') := r + \gamma \max_{a'} Q^{\mu'}(s, a' | \theta^{Q'})$$


$$II. \quad \theta^Q \leftarrow \theta^Q - \alpha \nabla_{\theta^Q} \left[ \left( Q^T(s, a, r, s') - Q^\mu(s, a | \theta^Q) \right)^2 \right]$$

iii. Every  $C$  steps reset  $\theta^{Q'} \leftarrow \theta^Q$

Target “Critic” Network ·  $Q^{\mu'}(s, a | \theta^{Q'})$  :



Replay Buffer ·  $\mathcal{R}$  :



# Algorithms Evolution

DQN · DPG · DDPG      *Deterministic Policy Gradient*


Expected Return:

$$J(\mu_\theta) := \mathbb{E}_{a_t = \mu_\theta(s_t)}[R(\tau)]$$


Deterministic Policy Gradient Theorem:

$$\nabla_{\theta^\mu} J(\mu_{\theta^\mu}) = \mathbb{E} \left[ \nabla_a Q^\mu(s, a \mid \theta^Q) \nabla_{\theta^\mu} \mu(s \mid \theta^\mu) \right]$$

“Critic” Network ·  $Q^\mu(s, a \mid \theta^Q)$ :



“Actor” Network ·  $\mu(s \mid \theta^\mu)$ :





# Algorithms Evolution


DQN · DPG · DDPG

Deep DPG


## Code Structure

1. Initialize randomly  $Q^\mu(s, a | \theta^Q)$  and  $\mu(s | \theta^\mu)$
2. Initialize  $Q^{\mu'}(s, a | \theta^{Q'} \leftarrow \theta^Q)$  and  $\mu'(s | \theta^{\mu'} \leftarrow \theta^\mu)$
3. Initialize replay buffer  $\mathcal{R} = S \times A \times R \times S$
4. Observe initial state  $s$
5. Repeat:
  - a.  · Sample
  - b.  · Train


“Critic” Network  
 $Q^\mu(s, a | \theta^Q)$




Target “Critic” Network  
 $Q^{\mu'}(s, a | \theta^{Q'})$




“Actor” Network  
 $\mu(s | \theta^\mu)$



Target “Actor” Network  
 $\mu'(s | \theta^{\mu'})$



Replay Buffer ·  $\mathcal{R}$



# Algorithms Evolution


DQN · DPG · DDPG

Deep DPG


■ · Sample:

- i. Execute  $a = \mu(s | \theta^\mu) + \mathcal{N}$  and observe  $r$  and  $s'$
- ii. Store  $\mathcal{R} \leftarrow \mathcal{R} \cup (s, a, r, s')$

“Critic” Network  
 $Q^\mu(s, a | \theta^Q)$



Target “Critic” Network  
 $Q^{\mu'}(s, a | \theta^{Q'})$




“Actor” Network  
 $\mu(s | \theta^\mu)$



Target “Actor” Network  
 $\mu'(s | \theta^{\mu'})$



Replay Buffer ·  $\mathcal{R}$



# Algorithms Evolution

DQN · DPG · DDPG

Deep DPG

■ · Train

i. Sample random batch  $\mathcal{B} \subseteq \mathcal{R}$

$$ii. Q^T(s, a, r, s') := r + \gamma Q^{\mu'}(s', \overbrace{\mu'(s' | \theta^{\mu'})}^{a'} | \theta^{Q'})$$

$$iii. L^Q(\theta^Q) = \frac{1}{|\mathcal{B}|} \sum_{(s,a,r,s') \in \mathcal{B}} [Q^T(s, a, r, s') - Q^\mu(s, a | \theta^Q)]^2$$

$$iv. \theta^Q \leftarrow \theta^Q - \alpha \nabla_{\theta^Q} L(\theta^Q)$$


$$v. L^\mu(\theta^\mu) = \frac{1}{|\mathcal{B}|} \sum_{(s,a,r,s') \in \mathcal{B}} \nabla_a Q^\mu(s, a | \theta^Q) \nabla_{\theta^\mu} \mu(s | \theta^\mu)$$

$$vi. \theta^\mu \leftarrow \theta^\mu - \alpha \nabla_{\theta^\mu} L(\theta^\mu)$$

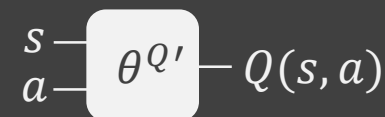
$$vii. \theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'} \text{ where } \tau \ll 1$$

$$viii. \theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$


“Critic” Network  
 $Q^\mu(s, a | \theta^Q)$




Target “Critic” Network  
 $Q^{\mu'}(s, a | \theta^{Q'})$




“Actor” Network  
 $\mu(s | \theta^\mu)$




Target “Actor” Network  
 $\mu'(s | \theta^{\mu'})$



Replay Buffer ·  $\mathcal{R}$



# DDPG Performance



· DPG

· DDPG

· DDPG + Batch Norm.


· DDPG + Batch Norm. + Raw pixels




# Conclusion

- + *Continuous control is much more stable.*
- + *Same algorithm/tuning solves many (20+) problems.*
- + *Pseudo-Code and hyperparameters are provided.*
- *Computationally expensive.*
- *Many hyperparameters must be tuned correctly.*
- *Many decisions of the paper require lot of pre-knowledge.*






DQN · Discrete



DDPG · Continuous



# Any Questions?



# References

## Presented Paper:

Lillicrap, Timothy P., et al. "Continuous control with deep reinforcement learning." *arXiv preprint arXiv:1509.02971* (2015).

## DQN Paper:

Mnih, Volodymyr, et al. "Playing atari with deep reinforcement learning." *arXiv preprint arXiv:1312.5602* (2013).

## DPG Paper:

Silver, David, et al. "Deterministic policy gradient algorithms." 2014.

## Useful Websites:

- [https://spinningup.openai.com/en/latest/spinningup/rl\\_intro.html](https://spinningup.openai.com/en/latest/spinningup/rl_intro.html)
- <https://towardsdatascience.com/deep-deterministic-policy-gradients-explained-2d94655a9b7b>
- [https://medium.com/@jonathan\\_hui/rl-dqn-deep-q-network-e207751f7ae4](https://medium.com/@jonathan_hui/rl-dqn-deep-q-network-e207751f7ae4)

# Hyperparameters

*Learning Rates*

$$\alpha_{actor} = 10^{-4}$$

$$\alpha_{critic} = 10^{-3}$$

*Discount Factor*

$$\gamma = 0.99$$

*Target update*

$$\tau = 0.001$$

*Neural network*

*ReLU for all except last layer  
of the actor that uses tanh, 2  
hidden layers of 300  
respectively 400 neurons.*

*Buffer, Batch Size*

$$10^6, 64$$

# Improvements after 2015

- *Paper uses batch normalization which has been found to improve only certain problems.*
- *An architecture with several actors can explore more of the environment.*